



COGNITIVE ANALYTICS LAB

Submitted by:

Wajid Ali Hashmi

500096923

B2(H)

Submitted to:

Sugandha Sharma

Index

Sno.	Name	Date	Page
1.	Analysing the dataset	17-01-2024	3
2.	Ridge and lasso algorithm	24-01-2024	6
3.	Logistic and Polynomial Regression	31-01-2024	11
4.	Implementing Decision Tree on diabetes dataset	14-02-2024	15
5.	Implementing Random Forest	21-02-2024	17
6.	Implementing Gradient Boosting	20-03-2024	20
7.	Implementing AdaBoost on Breast cancer dataset	25-03-2024	24
8.	Linear regression in R	10-04-2024	26
9.	K-means Clustering in R	10-04-2024	28

Experiment 1

Analysing the dataset based on some values

Dataset source: http://blog.hackerearth.com/wp-content/uploads/2016/12/airfoil_self_noise.csv

```
[17] import pandas as pd
import numpy as np
import statsmodels.api as sm
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

[18] df = pd.read_csv("D:\\DESKTOP\\cognitive analysis\\self noise.csv")
df.head()
```

	Frequency(Hz)	Angle of Attack	Chord Length	Free stream velocity	Displacement	Sound pressure level
0	800	0.0	0.3048	71.3	0.002663	126.201
1	1000	0.0	0.3048	71.3	0.002663	125.201
2	1250	0.0	0.3048	71.3	0.002663	125.951
3	1600	0.0	0.3048	71.3	0.002663	127.591
4	2000	0.0	0.3048	71.3	0.002663	127.461

```
[19] correlation_matrix=df.corr()
print(correlation_matrix)
```

	Frequency(Hz)	Angle of Attack	Chord Length	Free stream velocity	Displacement	Sound pressure level
Frequency(Hz)	1.000000	-0.272765	-0.003661	0.133664	-0.230107	-0.390711
Angle of Attack	-0.272765	1.000000	-0.504868	0.058760	0.753394	-0.156108
Chord Length	-0.003661	-0.504868	1.000000	0.003787	-0.220842	-0.236162
Free stream velocity	0.133664	0.058760	0.003787	1.000000	-0.003974	0.125103
Displacement	-0.230107	0.753394	-0.220842	-0.003974	1.000000	-0.312670
Sound pressure level	-0.390711	-0.156108	-0.236162	0.125103	-0.312670	1.000000

```
[20] X = df.drop(columns=['Sound_pressure_level'])
X=sm.add_constant(X)
y=df['Sound_pressure_level']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
Linear_Regression=sm.OLS(y_train, X_train).fit()

y_pred = Linear_Regression.predict(X_test)
r_squared = r2_score(y_test, y_pred)
adj_r_squared = Linear_Regression.rsquared_adj
f_statistic = Linear_Regression.fvalue
p_value = Linear_Regression.f_pvalue
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)
print(f"R-squared: {r_squared}")
print(f"Adjusted R-squared: {adj_r_squared}")
print(f"F-statistics: {f_statistic}, p-value: {p_value}")
print(f"MSE: {mse}")
print(f"RMSE: {rmse}")
print(f"MAE: {mae}")
```

```
[15] R-squared: 0.5582979754896911
Adjusted R-squared: 0.5013716488971148
F-statistics: 242.52150554359076, p-value: 6.080414968967961e-179
MSE: 22.128643318249154
RMSE: 4.704109194975087
MAE: 3.672414564174714
```


```

df_with_const = sm.add_constant(df[expected_features])
df['Predicted'] = LinearRegression.predict(df_with_const)
df['Residuals'] = df['Sound_pressure_level'] - df['Predicted']
plt.figure(figsize=(10, 6))
sns.scatterplot(x='Predicted', y='Residuals', data=df)
plt.title('Residual Plot')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.axhline(y=0, color='r', linestyle='--')
plt.show()

```

[27]

Python



Interpretations:

Based on the analysis, the following can be deduced about the dataset:

- The R-squared value of 0.558 indicates that the linear regression model explains 55.8% of the variance in the sound pressure level. This is a good fit for a linear model.
- The adjusted R-squared value of 0.501 is slightly lower than the R-squared value, but it still indicates a good fit for the model.
- The F-statistic of 242.52 and the p-value of less than 0.0001 both indicate that the model is statistically significant. This means that the coefficients of the model are not equal to zero.
- The MSE of 22.12 indicates that the average squared difference between the predicted and actual sound pressure levels is 22.12.
- The RMSE of 4.704 is the square root of the MSE, and it indicates that the average difference between the predicted and actual sound pressure levels is 4.704.
- The MAE of 3.672 indicates that the average absolute difference between the predicted and actual sound pressure levels is 3.672.
- The residual plot shows that the residuals are randomly scattered around the zero line. This indicates that the model does not have any heteroscedasticity or autocorrelation problems.

- The coefficients of the model can be used to identify which features have the most impact on the sound pressure level.
- The model can be used to make predictions about the sound pressure level for new data points.
- The model can be used to identify outliers in the data.

In conclusion, the linear regression model is a good fit for the data and can be used to predict the sound pressure level based on the other features in the dataset.

Experiment 2

Implementation of Ridge and lasso algorithm in Python:

Dataset source: <https://github.com/JWarmenhoven/ISLR-python/blob/master/Notebooks/Data/Hitters.csv>

```
[1] ✓ 18.1s Python
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import scale
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge, RidgeCV, Lasso, LassoCV
from sklearn.metrics import mean_squared_error

[2] ✓ 0.4s Python
df = pd.read_csv('Hitters.csv').dropna().drop('Player', axis = 1)
df.info()
dummies = pd.get_dummies(df[['League', 'Division', 'NewLeague']])

...
<class 'pandas.core.frame.DataFrame'>
Int64Index: 263 entries, 1 to 321
Data columns (total 20 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   AtBat       263 non-null     int64
1   Hits        263 non-null     int64
2   HmRun       263 non-null     int64
3   Runs        263 non-null     int64
4   RBI         263 non-null     int64
5   Walks       263 non-null     int64
6   Years       263 non-null     int64
7   CAtBat      263 non-null     int64
8   CHits       263 non-null     int64
9   CHmRun      263 non-null     int64
10  CRuns       263 non-null     int64
11  CRBI        263 non-null     int64
12  CWalks      263 non-null     int64
13  League      263 non-null     object
14  Division    263 non-null     object
15  PutOuts     263 non-null     int64
16  Assists     263 non-null     int64
17  Errors      263 non-null     int64
18  Salary      263 non-null     float64
19  NewLeague   263 non-null     object
dtypes: float64(1), int64(16), object(3)
memory usage: 43.1+ KB

[3] ✓ 0.0s Python
y = df.Salary
X_ = df.drop(['Salary', 'League', 'Division', 'NewLeague'], axis = 1).astype('float64')
# Define the feature set X.
X = pd.concat([X_, dummies[['League_N', 'Division_W', 'NewLeague_N']]], axis = 1)
X.info()

...
<class 'pandas.core.frame.DataFrame'>
Int64Index: 263 entries, 1 to 321
Data columns (total 19 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   AtBat       263 non-null     float64
1   Hits        263 non-null     float64
2   HmRun       263 non-null     float64
3   Runs        263 non-null     float64
4   RBI         263 non-null     float64
5   Walks       263 non-null     float64
6   Years       263 non-null     float64
7   CAtBat      263 non-null     float64
8   CHits       263 non-null     float64
9   CHmRun      263 non-null     float64
10  CRuns       263 non-null     float64
11  CRBI        263 non-null     float64
12  CWalks      263 non-null     float64
13  PutOuts     263 non-null     float64
14  Assists     263 non-null     float64
15  Errors      263 non-null     float64
16  League_N    263 non-null     uint8
17  Division_W  263 non-null     uint8
18  NewLeague_N 263 non-null     uint8
dtypes: float64(16), uint8(3)
memory usage: 35.7 KB
```

```

alphas = 10**np.linspace(10,-2,100)*0.5
alphas
[4] ✓ 0.0s Python
... array([5.00000000e+09, 3.78231664e+09, 2.86118383e+09, 2.16438064e+09,
1.63727458e+09, 1.23853818e+09, 9.36908711e+08, 7.08737081e+08,
5.36133611e+08, 4.05565415e+08, 3.06795364e+08, 2.32079442e+08,
1.75559587e+08, 1.32804389e+08, 1.00461650e+08, 7.59955541e+07,
5.74878498e+07, 4.34874501e+07, 3.28966612e+07, 2.48851178e+07,
1.88246790e+07, 1.42401793e+07, 1.07721735e+07, 8.14875417e+06,
6.16423370e+06, 4.66301673e+06, 3.52740116e+06, 2.66834962e+06,
2.01850863e+06, 1.52692775e+06, 1.15506485e+06, 8.73764200e+05,
6.60970574e+05, 5.00000000e+05, 3.78231664e+05, 2.86118383e+05,
2.16438064e+05, 1.63727458e+05, 1.23853818e+05, 9.36908711e+04,
7.08737081e+04, 5.36133611e+04, 4.05565415e+04, 3.06795364e+04,
2.32079442e+04, 1.75559587e+04, 1.32804389e+04, 1.00461650e+04,
7.59955541e+03, 5.74878498e+03, 4.34874501e+03, 3.28966612e+03,
2.48851178e+03, 1.88246790e+03, 1.42401793e+03, 1.07721735e+03,
8.14875417e+02, 6.16423370e+02, 4.66301673e+02, 3.52740116e+02,
2.66834962e+02, 2.01850863e+02, 1.52692775e+02, 1.15506485e+02,
8.73764200e+01, 6.60970574e+01, 5.00000000e+01, 3.78231664e+01,
2.86118383e+01, 2.16438064e+01, 1.63727458e+01, 1.23853818e+01,
9.36908711e+00, 7.08737081e+00, 5.36133611e+00, 4.05565415e+00,
3.06795364e+00, 2.32079442e+00, 1.75559587e+00, 1.32804389e+00,
1.00461650e+00, 7.59955541e-01, 5.74878498e-01, 4.34874501e-01,
3.28966612e-01, 2.48851178e-01, 1.88246790e-01, 1.42401793e-01,
1.07721735e-01, 8.14875417e-02, 6.16423370e-02, 4.66301673e-02,
3.52740116e-02, 2.66834962e-02, 2.01850863e-02, 1.52692775e-02,
1.15506485e-02, 8.73764200e-03, 6.60970574e-03, 5.00000000e-03])

```

```

ridge = Ridge(normalize = True)
coefs = []
for a in alphas:
    ridge.set_params(alpha = a)
    ridge.fit(X, y)
    coefs.append(ridge.coef_)
np.shape(coefs)
[5] ✓ 0.3s Python

```

```

... (100, 19)

```

```


ax = plt.gca()
ax.plot(alphas, coefs)
ax.set_xscale('log')
plt.axis('tight')
plt.xlabel('alpha')
plt.ylabel('weights')
[6] ✓ 1.0s Python

```

```

... Text(0, 0.5, 'weights')

```



```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.5, random_state=1)

```

```

ridge2 = Ridge(alpha = 4, normalize = True)
ridge2.fit(X_train, y_train) # Fit a ridge regression on the training data
pred2 = ridge2.predict(X_test) # Use this model to predict the test data
print(pd.Series(ridge2.coef_, index = X.columns)) # Print coefficients
print(mean_squared_error(y_test, pred2)) # Calculate the test MSE

```

```

ridge3 = Ridge(alpha = 10**10, normalize = True)
ridge3.fit(X_train, y_train) # Fit a ridge regression on the training data
pred3 = ridge3.predict(X_test) # Use this model to predict the test data
print(pd.Series(ridge3.coef_, index = X.columns)) # Print coefficients
print(mean_squared_error(y_test, pred3)) # Calculate the test MSE

```

```

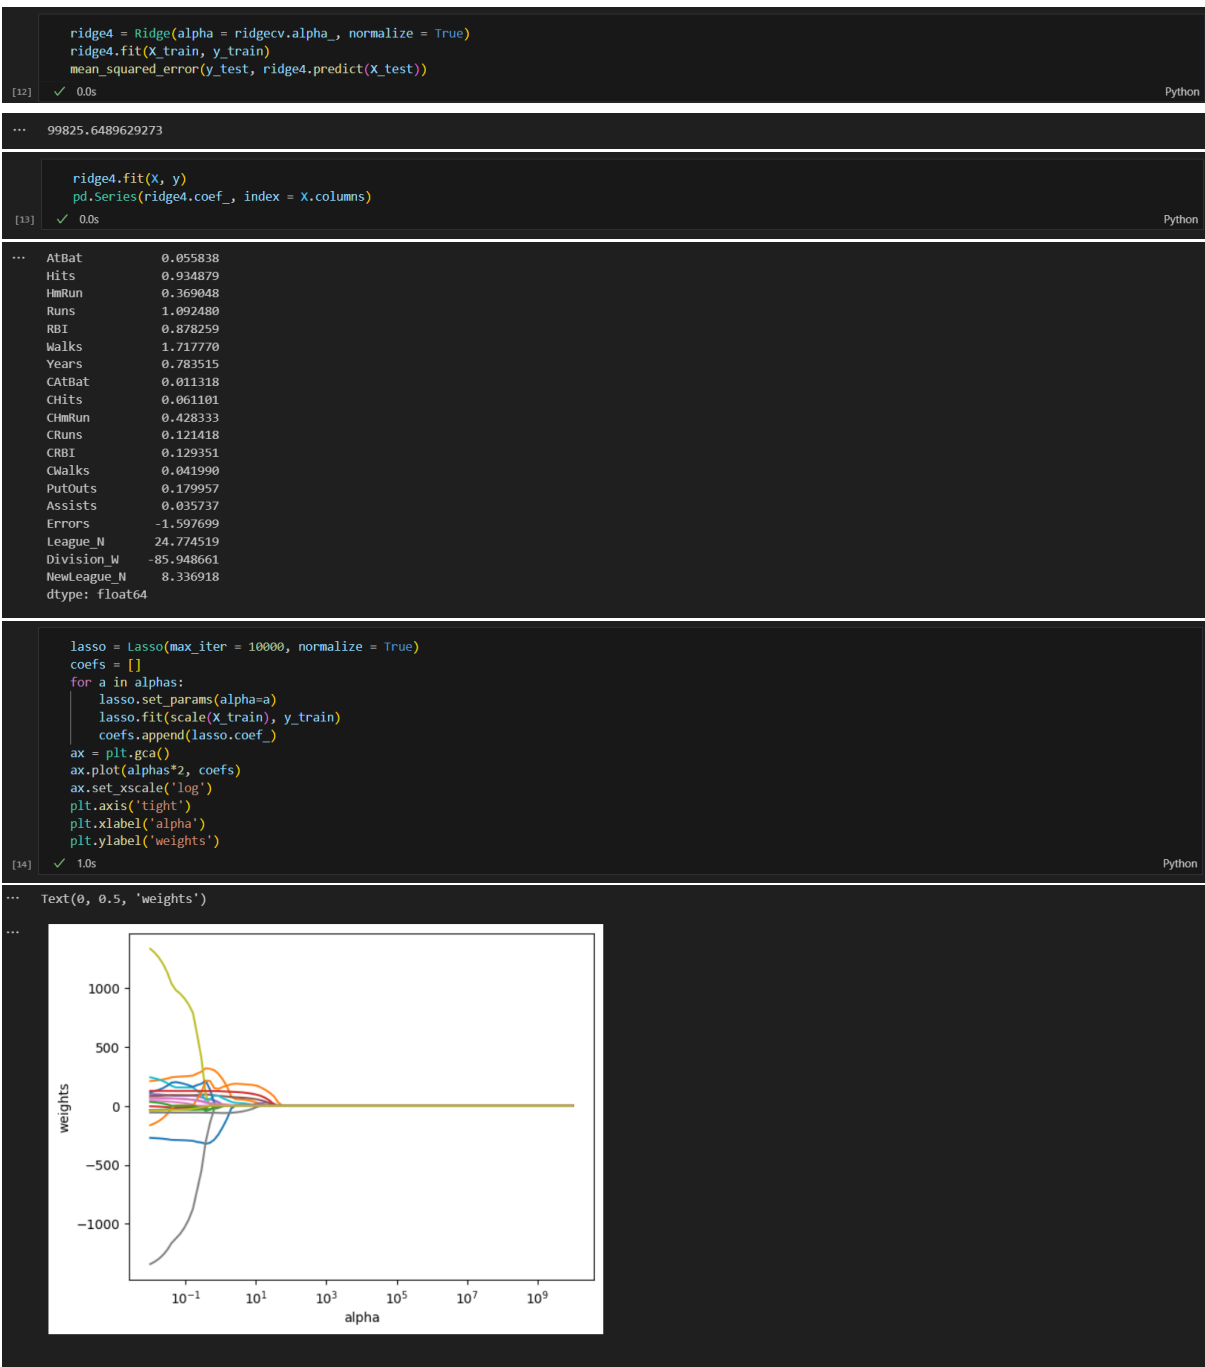
ridge2 = Ridge(alpha = 0, normalize = True)
ridge2.fit(X_train, y_train) # Fit a ridge regression on the training data
pred = ridge2.predict(X_test) # Use this model to predict the test data
print(pd.Series(ridge2.coef_, index = X.columns)) # Print coefficients
print(mean_squared_error(y_test, pred)) # Calculate the test MSE

```

```

ridgecv = RidgeCV(alphas = alphas, scoring = 'neg_mean_squared_error', normalize = True)
ridgecv.fit(X_train, y_train)
ridgecv.alpha_

```

```

lassocv = LassoCV(alphas = None, cv = 10, max_iter = 100000, normalize = True)
lassocv.fit(X_train, y_train)
lasso.set_params(alpha=lassocv.alpha_)
lasso.fit(X_train, y_train)
mean_squared_error(y_test, lasso.predict(X_test))

```

[15] ✓ 0.5s Python

... 104960.65853895503

```

pd.Series(lasso.coef_, index=X.columns)

```

[16] ✓ 0.0s Python

...

AtBat	0.000000
Hits	1.082446
HmRun	0.000000
Runs	0.000000
RBI	0.000000
Walks	2.906388
Years	0.000000
CAtBat	0.000000
CHits	0.000000
CHmRun	0.219367
CRuns	0.000000
CRBI	0.513975
CWalks	0.000000
PutOuts	0.368401
Assists	-0.000000
Errors	-0.000000
League_N	0.000000
Division_W	-89.064338
NewLeague_N	0.000000

dtype: float64

Interpretation:

- Ridge and lasso regression are regularization techniques used to prevent overfitting in linear regression models.
- The choice of alpha determines the strength of regularization, with higher alpha leading to more regularization.
 - As we can see from the output of the code, when the value of alpha is small the output values are significantly small but when the alpha is increased the output values decreases with higher rate.
- The code assesses model performance using mean squared error on the test set for both ridge and lasso regression.
- Cross-validation is employed to find optimal alpha values for ridge and lasso, enhancing model generalization.

Experiment 3

Logistic Regression:

Dataset URL: <https://www.kaggle.com/datasets/gauravtopre/bank-customer-churn-dataset>

```
#Logistic regression
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

[1] ✓ 6.1s Python

df = pd.read_csv('Bank Customer Churn Prediction.csv')

[2] ✓ 0.4s Python

X = df[['credit_score', 'age', 'tenure', 'balance', 'products_number', 'credit_card', 'active_member', 'estimated_salary', 'gender', 'country']]
X = pd.get_dummies(X, columns=['gender', 'country'], drop_first=True)
y = df['churn']

[3] ✓ 0.0s Python

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

[4] ✓ 0.0s Python

model = LogisticRegression(random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)

[5] ✓ 0.0s Python

print(f"Accuracy: {accuracy}")
print("\nConfusion Matrix:")
print(conf_matrix)
print("\nClassification Report:")
print(classification_rep)
#TN FP
#FN TP

[6] ✓ 0.0s Python

... Accuracy: 0.811

Confusion Matrix:
[[1543  64]
 [ 314  79]]

Classification Report:
      precision    recall  f1-score   support

0         0.83       0.96       0.89       1607
1         0.55       0.20       0.29        393

accuracy         0.69
macro avg        0.69       0.58       0.59       2000
weighted avg     0.78       0.81       0.77       2000
```

Interpretations on the basis of the values:

Accuracy:

Accuracy is the ratio of correctly predicted instances to the total instances. It is a general measure of the model's overall correctness. In this case, an accuracy of 0.811 (or 81.1%) means that the model correctly predicted the churn status for approximately 81.1% of the instances in the dataset.

Confusion Matrix:

A confusion matrix is a table that describes the performance of a classification model. It consists of four terms: True Positive (TP), True Negative (TN), False Positive (FP), and False Negative (FN). In your case:

- True Positive (TP): 79 instances where the model correctly predicted churn (1).
- True Negative (TN): 1543 instances where the model correctly predicted no churn (0).
- False Positive (FP): 64 instances where the model incorrectly predicted churn when there was no churn.
- False Negative (FN): 314 instances where the model incorrectly predicted no churn when there was churn.

Classification Report:

Precision, recall, and F1-score are metrics that provide more insights, especially in imbalanced datasets.

- Precision:
 - Precision is the ratio of correctly predicted positive observations to the total predicted positives.
 - $\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$
 - In this case, the precision for class 1 (churn) is approximately 0.55, indicating that 55% of the instances predicted as churn were actually churn.
- Recall (Sensitivity):
 - Recall is the ratio of correctly predicted positive observations to the total actual positives.
 - $\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$
 - In this case, the recall for class 1 (churn) is approximately 0.20, indicating that the model correctly identified 20% of the actual churn instances.
- F1-score:
 - F1-score is the weighted average of precision and recall.
 - $\text{F1-score} = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$

It balances precision and recall, providing a single metric that considers both false positives and false negatives.

The weighted average of the F1-scores for both classes is given in the report.

- Support:
 - Support is the number of actual occurrences of each class in the specified dataset.
- Macro Average and Weighted Average:
 - Macro average calculates the unweighted average of precision, recall, and F1-score for all classes.
 - Weighted average considers the number of instances for each class, providing more weight to the majority class.

Polynomial Regression:

```
#Polynomial regression
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import numpy as np

[7] ✓ 0.0s Python

X = df[['credit_score', 'age', 'tenure', 'balance', 'products_number', 'credit_card', 'active_member', 'estimated_salary', 'gender', 'country']]
y = df['churn']
X = pd.get_dummies(X, columns=['gender', 'country'], drop_first=True)

[8] ✓ 0.0s Python

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
poly = PolynomialFeatures(degree=2)
X_train_poly = poly.fit_transform(X_train)
X_test_poly = poly.transform(X_test)

[9] ✓ 0.0s Python

poly_reg = LinearRegression()
poly_reg.fit(X_train_poly, y_train)
y_pred_poly = poly_reg.predict(X_test_poly)
mse = mean_squared_error(y_test, y_pred_poly)
rmse = np.sqrt(mse)
mae = mean_absolute_error(y_test, y_pred_poly)
r2 = r2_score(y_test, y_pred_poly)

[10] ✓ 0.0s Python

print(f"Mean Squared Error (MSE): {mse}")
print(f"Root Mean Squared Error (RMSE): {rmse}")
print(f"Mean Absolute Error (MAE): {mae}")
print(f"R-squared (R2): {r2}")

[11] ✓ 0.0s Python

... Mean Squared Error (MSE): 0.11198360191185681
Root Mean Squared Error (RMSE): 0.33463951038670975
Mean Absolute Error (MAE): 0.2396460943992281
R-squared (R2): 0.2907391364316939
```

- Mean Squared Error (MSE):
 - MSE measures the average squared difference between the actual and predicted values.
 - A lower MSE indicates better model performance.
 - In this case, an MSE of approximately 0.112 means, on average, the squared difference between the actual and predicted churn values is 0.112.

- Root Mean Squared Error (RMSE):
 - RMSE is the square root of the MSE and provides a measure of the average absolute error.
 - It is expressed in the same units as the target variable.
 - In this case, an RMSE of approximately 0.335 means, on average, the absolute difference between the actual and predicted churn values is around 0.335.
- Mean Absolute Error (MAE):
 - MAE is the average absolute difference between the actual and predicted values.
 - It is less sensitive to outliers compared to MSE.
 - In this case, an MAE of approximately 0.240 means, on average, the absolute difference between the actual and predicted churn values is around 0.240.
- R-squared (R²):
 - R-squared measures the proportion of the variance in the dependent variable (churn) that is predictable from the independent variables (features).
 - R-squared values range from 0 to 1, where 1 indicates a perfect fit.
 - In this case, an R-squared of approximately 0.291 means that around 29.1% of the variance in churn can be explained by the features in the model.

Experiment 4

Implementing Decision Tree on diabetes dataset:

Dataset URL: <https://www.kaggle.com/datasets/uciml/pima-indians-diabetes-database>

```
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn import metrics

col_names = ['pregnant', 'glucose', 'bp', 'skin', 'insulin', 'bmi', 'pedigree', 'age', 'label']
df = pd.read_csv("diabetes.csv", header=None, names=col_names)
df.drop(0, inplace=True)
df.head()
```

	pregnant	glucose	bp	skin	insulin	bmi	pedigree	age	label
1	6	148	72	35	0	33.6	0.627	50	1
2	1	85	66	29	0	26.6	0.351	31	0
3	8	183	64	0	0	23.3	0.672	32	1
4	1	89	66	23	94	28.1	0.167	21	0
5	0	137	40	35	168	43.1	2.288	33	1

```
feature_cols = ['pregnant', 'insulin', 'bmi', 'age', 'glucose', 'bp', 'pedigree']
X = df[feature_cols]
y = df.label

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)

d_tree = DecisionTreeClassifier()
d_tree = d_tree.fit(X_train, y_train)
y_pred = d_tree.predict(X_test)

print("Accuracy:", metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7056277056277056

Interpretations:

In this code we can see the following things:

- The value 0.7056277056277056 represents the accuracy score, which is a metric used to evaluate the performance of a machine learning model.
- The higher the accuracy score, the better the model is at making correct predictions.

Enhancing the accuracy of the decision tree:



```
clf = DecisionTreeClassifier(criterion="entropy", max_depth=3)
clf = clf.fit(X_train,y_train)
y_pred = clf.predict(X_test)

print("Accuracy:",metrics.accuracy_score(y_test, y_pred))
```

Accuracy: 0.7705627705627706

- Here we can see that the accuracy of the model increases when we ass the criterion of ‘entropy’ and ‘max_depth’
- The higher the accuracy, the better the model is at making predictions.

Experiment 5

Implementing Random Forest:

Dataset URL: <https://www.kaggle.com/datasets/yufengsui/portuguese-bank-marketing-data-set>

```
[49] import pandas as pd
import numpy as np
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, precision_score, recall_score, ConfusionMatrixDisplay
from sklearn.model_selection import RandomizedSearchCV, train_test_split
from scipy.stats import randint
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import export_graphviz
from IPython.display import Image
from sklearn import tree
from matplotlib import pyplot as plt
import graphviz

Python
```

```
[50] df = pd.read_csv('bank_cleaned.csv')
df.info()

Python
```

```
... <class 'pandas.core.frame.DataFrame'>
RangeIndex: 40841 entries, 0 to 40840
Data columns (total 18 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Unnamed: 0   40841 non-null  int64
1   age          40841 non-null  int64
2   job          40841 non-null  object
3   marital      40841 non-null  object
4   education    40841 non-null  object
5   default      40841 non-null  object
6   balance      40841 non-null  int64
7   housing      40841 non-null  object
8   loan         40841 non-null  object
9   day          40841 non-null  int64
10  month        40841 non-null  object
11  duration     40841 non-null  float64
12  campaign     40841 non-null  int64
13  pdays        40841 non-null  int64
14  previous     40841 non-null  int64
15  poutcome     40841 non-null  object
16  response     40841 non-null  object
17  response_binary 40841 non-null  int64
dtypes: float64(1), int64(8), object(9)
memory usage: 5.6+ MB
```

```
[51] df.head()

Python
```

```
...
   Unnamed: 0  age  job   marital  education  default  balance  housing  loan  day  month  duration  campaign  pdays  previous  poutcome  resp
0           0   58  management  married    tertiary    no     2143     yes    no   5   may       4.35         1     -1         0   unknown
1           1   44  technician  single    secondary    no        29     yes    no   5   may       2.52         1     -1         0   unknown
2           2   33  entrepreneur  married    secondary    no         2     yes   yes   5   may       1.27         1     -1         0   unknown
3           5   35  management  married    tertiary    no      231     yes    no   5   may       2.32         1     -1         0   unknown
4           6   28  management  single     tertiary    no      447     yes   yes   5   may       3.62         1     -1         0   unknown
```

```
[52] X = df.drop(['response', 'response_binary'], axis=1)
y = df['response_binary']
Enc= LabelEncoder()
for column in X.select_dtypes(include = ['object']).columns:
    X[column] = Enc.fit_transform(X[column])

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

Python
```

```

rf = RandomForestClassifier()
rf.fit(X_train, y_train)
Python

RandomForestClassifier
RandomForestClassifier()

y_pred = rf.predict(X_test)
Python

accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
Python
Accuracy: 0.9135757130615743

param_dist = {'n_estimators': randint(50,500), 'max_depth': randint(1,20)}
rf = RandomForestClassifier()
Python

rand_search = RandomizedSearchCV(rf, param_distributions = param_dist, n_iter=5, cv=5)
rand_search.fit(X_train, y_train)
Python

RandomizedSearchCV
estimator: RandomForestClassifier
RandomForestClassifier

best_rf = rand_search.best_estimator_
print('Best hyperparameters:', rand_search.best_params_)
Python
Best hyperparameters: {'max_depth': 12, 'n_estimators': 329}

cm = confusion_matrix(y_test, y_pred)
ConfusionMatrixDisplay(confusion_matrix=cm).plot()
Python
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x1a4b3caab40>

True label
0 6963 203
1 525 478
Predicted label 0 1
6000
5000
4000
3000
2000
1000

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
print("Accuracy:", accuracy)
print("Precision:", precision)
print("Recall:", recall)
Python
Accuracy: 0.9108826049700086
Precision: 0.7019089574155654
Recall: 0.4765702891326022

```

Interpretation over the metrics calculated:

The confusion matrix shows the performance of a random forest model on a classification task. The rows represent the actual labels, and the columns represent the predicted labels. Each cell shows the number of instances in a particular category. In this case, the model is trying to predict whether a customer will churn or not (represented by 0 and 1).

1. Accuracy (0.91): This metric tells us how well the model performed overall. An accuracy of 91% means that the model classified 91% of the instances correctly.
2. Precision (0.70): This metric tells us the proportion of positive predictions that were actually correct. A precision of 70% means that out of all the instances that the model predicted to churn, 70% of them actually churned.
3. Recall (0.48): This metric tells us the proportion of actual positive instances that were identified correctly. A recall of 48% means that out of all the customers who actually churned, the model identified only 48% of them.

In simpler terms, the model is very good at identifying non-churning customers (high accuracy), but it struggles to identify churning customers (low recall). This could be because churning customers are a smaller proportion of the overall dataset, so the model is less likely to learn their characteristics.

Experiment 6

Implementing Gradient Boosting:

Dataset: <https://www.kaggle.com/lodetomasi1995/income-classification>

```
from sklearn.ensemble import GradientBoostingClassifier
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn import preprocessing
#from sklearn import warnings.filterwarnings("ignore")

C:\Users\Wajid\AppData\Local\Temp\ipykernel_32736\529460093.py:3: DeprecationWarning:
Pyarrow will become a required dependency of pandas in the next major release of pandas (pandas 3.0),
(to allow more performant data types, such as the Arrow string type, and better interoperability with other libraries)
but was not found to be installed on your system.
If this would cause problems for you,
please provide us feedback at https://github.com/pandas-dev/pandas/issues/54466

import pandas as pd

df = pd.read_csv('D:\DESKTOP\Cognitive Analytics\income classification\income_evaluation.csv')
df.head()
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	income
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

```
df.columns

Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',
       'marital-status', 'occupation', 'relationship', 'race', 'sex',
       'capital-gain', 'capital-loss', 'hours-per-week', 'native-country',
       'income'],
      dtype='object')

df.columns = df.columns.str.strip()
X_train, X_test, y_train, y_test = train_test_split(df.drop('income', axis=1),
                                                    df['income'], test_size=0.2)

gbc = GradientBoostingClassifier(n_estimators=500, learning_rate=0.05, random_state=100, max_features=5 )

X_train_encoded = pd.get_dummies(X_train)
X_test_encoded = pd.get_dummies(X_test)
gbc.fit(X_train_encoded, y_train)

GradientBoostingClassifier(learning_rate=0.05, max_features=5, n_estimators=500,
                           random_state=100)

X_train_encoded = pd.get_dummies(X_train)
X_test_encoded = pd.get_dummies(X_test)
X_test_encoded = X_test_encoded.reindex(columns=X_train_encoded.columns, fill_value=0)
print(confusion_matrix(y_test, gbc.predict(X_test_encoded)))

[[4689 245]
 [ 626 953]]
```

```

print("GBC accuracy is %2.2f" % accuracy_score(y_test, gbc.predict(X_test_encoded)))
[8] Python
... GBC accuracy is 0.87

from sklearn.metrics import classification_report
pred = gbc.predict(X_test_encoded)
print(classification_report(y_test, pred))
[9] Python
...
      precision    recall  f1-score   support

<=50K      0.88       0.95       0.92       4934
>50K       0.80       0.60       0.69       1579

accuracy          0.84       0.78       0.80       6513
macro avg         0.86       0.87       0.86       6513
weighted avg      0.86       0.87       0.86       6513

from sklearn.model_selection import GridSearchCV
[10] Python

grid = { 'learning_rate':[0.01,0.05,0.1],
        'n_estimators':np.arange(100,500,100), }
[11] Python

gb = GradientBoostingClassifier()
gb_cv = GridSearchCV(gb, grid, cv = 4)
gb_cv.fit(X_train_encoded, y_train)
print("Fitting completed successfully")
[12] Python
... Fitting completed successfully

>
#Print("Is the GridSearchCV fitted?", gb_cv.cv_results_ is not None)
print("Available attributes:", dir(gb_cv))
# Check if the object is correctly named and assigned
print(gb_cv)

#Print("Best Estimator:", gb_cv.best_estimator_)
print("'cv_results_' available:", hasattr(gb_cv, "cv_results_"))
print("'best_params_' available:", hasattr(gb_cv, "best_params_"))
print("cv_results_:", gb_cv.best_index_)
[13] Python
... Available attributes: ['__abstractmethods__', '__annotations__', '__class__', '__delattr__', '__dict__', '__dir__', '__doc__', '__eq__', '__format__', '__ge__',
GridSearchCV(cv=4, estimator=GradientBoostingClassifier(),
param_grid={'learning_rate': [0.01, 0.05, 0.1],
'n_estimators': array([100, 200, 300, 400])})
'cv_results_' available: True
'best_params_' available: True
cv_results_: 11

print("Best Parameters:",gb_cv.best_params_)
print("Train Score:",gb_cv.best_score_)
print("Test Score:",gb_cv.score(X_test_encoded,y_test))
[14] Python
... Best Parameters: {'learning_rate': 0.1, 'n_estimators': 400}
Train Score: 0.871237714987715
Test Score: 0.8756333486872409

>
grid = {'max_depth':[2,3,4,5,6,7] }
gb = GradientBoostingClassifier(learning_rate=0.1,n_estimators=400)
gb_cv = GridSearchCV(gb, grid, cv = 4)
gb_cv.fit(X_train_encoded,y_train)
[15] Python

> GridSearchCV ⓘ ⓘ
> estimator: GradientBoostingClassifier
  GradientBoostingClassifier ⓘ

>
print("Best Parameters:",gb_cv.best_params_)
print("Train Score:",gb_cv.best_score_)
print("Test Score:",gb_cv.score(X_test_encoded,y_test))
[16] Python
... Best Parameters: {'max_depth': 3}
Train Score: 0.8711609336609336
Test Score: 0.8756333486872409

```

Interpretation of the Confusion Matrix and Model Performance

Based on the confusion matrix and performance metrics, here's an interpretation of the model's behaviour on the dataset:

Confusion Matrix:

- **Correct Classifications:**
 - The model correctly classified 4689 data points in the $\leq 50K$ class and 953 data points in the $> 50K$ class.
 - This indicates a good ability to identify data points belonging to each class.
- **Incorrect Classifications:**
 - The model misclassified 245 data points from $\leq 50K$ class and 626 data points from the $> 50K$ class.
 - There's a higher number of misclassifications for the $> 50K$ class, suggesting the model struggles more with this category.

Performance Metrics:

- **Overall Accuracy: 87%** - This indicates a good overall performance of the model in classifying data points correctly.
- **Precision:**
 - $\leq 50K$: 88% - The model is good at identifying true positives for the $\leq 50K$ class (out of all predicted positives, 88% are actually positive).
 - $> 50K$: 80% - The model is less precise for the $> 50K$ class, meaning there might be more false positives (predicted positive but actual negative).

Recall:

- $\leq 50K$: 95% - The model captures most of the actual positives in the $\leq 50K$ class (out of all actual positives, 95% are predicted positive).
- $> 50K$: 60% - The model misses a significant portion (40%) of the actual positives in the $> 50K$ class (false negatives). This is the major contributor to the lower precision for this class.

F1-Score: This metric balances precision and recall. It's higher for the $\leq 50K$ class (0.92) compared to the $> 50K$ class (0.69), again reflecting the model's better performance for the former.

Overall Interpretation:

The model performs well in classifying data points with an overall accuracy of 87%. It excels at identifying $\leq 50K$ class data points with high precision (88%) and recall (95%). However, the model struggles with the $>50K$ class, particularly with recall (60%), leading to more false negatives (missed positives). This suggests the model might be biased towards the majority class ($\leq 50K$) or have difficulty learning the patterns in the $>50K$ class data.

Experiment 7

AdaBoost Classifier on Breast cancer dataset:

```
[11] ✓ 0.0s Python
1 import pandas as pd
2 import numpy as np
3 from sklearn.ensemble import AdaBoostClassifier
4 from sklearn.tree import DecisionTreeClassifier
5 from sklearn.datasets import load_breast_cancer
6 from sklearn.model_selection import train_test_split
7 from sklearn.metrics import confusion_matrix, accuracy_score
8 from sklearn.preprocessing import LabelEncoder
9 import warnings
10 warnings.filterwarnings("ignore")

[12] ✓ 0.0s Python
1 breast_cancer = load_breast_cancer()

[13] ✓ 0.0s Python
• 1 .DataFrame(breast_cancer.data, columns = breast_cancer.feature_names)
  2 .Categorical.from_codes(breast_cancer.target, breast_cancer.target_names)

[14] ✓ 0.0s Python
1 encoder = LabelEncoder()
2 binary_encoded_y = pd.Series(encoder.fit_transform(y))

[15] ✓ 0.0s Python
1 train_X, test_X, train_y, test_y = train_test_split(X, binary_encoded_y, random_state = 1)

[16] ✓ 0.5s Python
1 classifier = AdaBoostClassifier()
2 classifier.fit(train_X, train_y)
3 prediction = classifier.predict(test_X)

[17] ✓ 0.0s Python
1 confusion_matrix(test_y, prediction)
... array([[85,  3],
          [ 3, 52]], dtype=int64)

[18] ✓ 0.0s Python
1 accuracy = accuracy_score(test_y, prediction)
2 print('AdaBoost Accuracy: ', accuracy)
... AdaBoost Accuracy:  0.958041958041958
```


Interpretation:

Based on the confusion matrix and accuracy, your AdaBoost classifier seems to be performing very well. Here's a breakdown:

- High True Positives (85) and True Negatives (52): This indicates the classifier effectively identified most of the actual positive and negative instances.
- Low False Positives (3) and False Negatives (3): The low number of misclassified instances suggests the classifier generalizes well and doesn't overfit to the training data.

Overall, your AdaBoost classifier demonstrates a high level of accuracy in distinguishing between the positive and negative classes based on the provided confusion matrix and accuracy.

Experiment 8

Linear Regression in R:

Dataset: <https://ocw.mit.edu/courses/15-097-prediction-machine-learning-and-statistics-spring-2012/resources/salary/>

```
dataset <- read.csv("D:/DESKTOP/Cognitive Analytics/Experiments/Datasets/Salary.csv")
library(caTools)
split <- sample.split(dataset$Salary, SplitRatio = 0.7)
trainingset <- subset(dataset, split == TRUE)
testset <- subset(dataset, split == FALSE)
lm.r <- lm(formula = Salary ~ YearsExperience, data = trainingset)
coef(lm.r)
ypred <- predict(lm.r, newdata = testset)
library(ggplot2)
ggplot() + geom_point(aes(x = trainingset$YearsExperience,
                           y = trainingset$Salary), colour = 'red')
+ geom_line(aes(x = trainingset$YearsExperience, y = predict(lm.r, newdata = trainingset)),
            colour = 'blue') +
  ggtitle('Salary vs Experience (Training set)') +
  xlab('Years of experience') +
  ylab('Salary')

ggplot() + geom_point(aes(x = testset$YearsExperience,
                           y = testset$Salary),
                      colour = 'red') +
  geom_line(aes(x = trainingset$YearsExperience,
                 y = predict(lm.r, newdata = trainingset)),
            colour = 'blue') +
  ggtitle('Salary vs Experience (Test set)') +
  xlab('Years of experience') +
  ylab('Salary')
```





Interpretations:

Intercept (30732.10): This value represents the expected value of the dependent variable when the independent variable (YearsExperience) is equal to zero. In simpler terms, it's the point where the regression line crosses the y-axis.

YearsExperience (8557.83): This coefficient represents the change in the dependent variable for every one unit increase in the independent variable (YearsExperience). It signifies the slope of the regression line. Since the coefficient is positive (8557.83), we can interpret that there's a positive linear relationship between years of experience and the dependent variable. In other words, as the years of experience increase by one unit, the dependent variable is expected to increase by 8557.83 units on average.

Experiment 9

K-means Clustering in R:

Dataset: <https://www.kaggle.com/datasets/saurabh00007/iriscsv>

```
library(clusterR)
library(cluster)

data(iris)
str(iris)

iris_1 <- iris[, -5]
set.seed(240)


kmeans.re <- kmeans(iris_1, centers = 3, nstart = 20)
kmeans.re
kmeans.re$cluster


cm <- table(iris$Species, kmeans.re$cluster)
cm

plot(iris_1[c("Sepal.Length", "Sepal.Width")])
plot(iris_1[c("Sepal.Length", "Sepal.Width")], col = kmeans.re$cluster)


plot(iris_1[c("Sepal.Length", "Sepal.Width")], col = kmeans.re$cluster, main = "K-means with 3 clusters")
kmeans.re$centers
kmeans.re$centers[, c("Sepal.Length", "Sepal.Width")]
points(kmeans.re$centers[, c("Sepal.Length", "Sepal.Width")], col = 1:3, pch = 8, cex = 3)


y_kmeans <- kmeans.re$cluster
clusplot(iris_1[, c("Sepal.Length", "Sepal.Width")],
  y_kmeans,
  lines = 0,
  shade = TRUE,
  color = TRUE,
  labels = 2,
  plotchar = FALSE,
  span = TRUE,
  main = paste("Cluster iris"),
  xlab = 'Sepal.Length',
  ylab = 'Sepal.Width')
```





K-means with 3 clusters





These two components explain 100 % of the point variability.

Results and Interpretation:

The output (kmeans.re) shows that the algorithm successfully assigned data points to three clusters with sizes 50, 62, and 38.

The kmeans.re\$centers table displays the average values (centroids) for each cluster in terms of sepal and petal length and width.

- Cluster 1 (centroid: Sepal Length - 5.01, Sepal Width - 3.43) seems to represent flowers with the smallest overall dimensions.
- Cluster 2 (centroid: Sepal Length - 5.90, Sepal Width - 2.75) appears to group flowers with slightly larger sepals but narrower sepals compared to Cluster 3.
- Cluster 3 (centroid: Sepal Length - 6.85, Sepal Width - 3.07) likely represents flowers with the largest sepals and wider sepals on average.

Visualized the cluster assignments (`kmeans.re$cluster`) in two ways:

- By colouring data points in a scatter plot based on their assigned cluster. This visually highlights the separation between clusters in the two-dimensional space of sepal length and width.
- By using the `clusplot` function from the `ClusterR` library, which creates a more comprehensive cluster plot with shading and labels.